

# Cloud Native Fundamentals

## 2.1 Introduction to Cloud Native Fundamentals

### Cores

1. architecture considerations
2. container orchestration
3. open source PaaS
4. cloud native CI/CD

### Cloud native

- build at scale
- on-premise, hybrid or public cloud
- speed + agility

### Containers

- small unit of applicat
- deploy fast, manageable
- fits with microservice-based architecture

### Container Orchestration - manage containers

- Docker Swarm
- Apache Mesos
- Kubernetes - automates - ~~automate~~  
configuration  
management  
scalability

- runtime
- network
- storage
- service mesh
- logs + metrics
- tracing

## 2.2 Architecture Considerations

### Cover

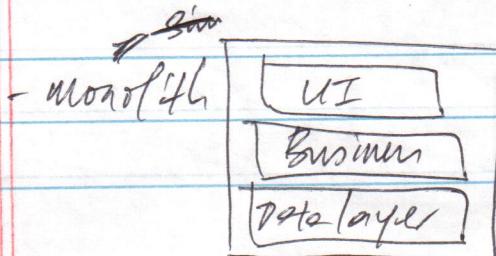
1. Monolith + microservices
2. Trade-offs
3. Practices for application development

### Design

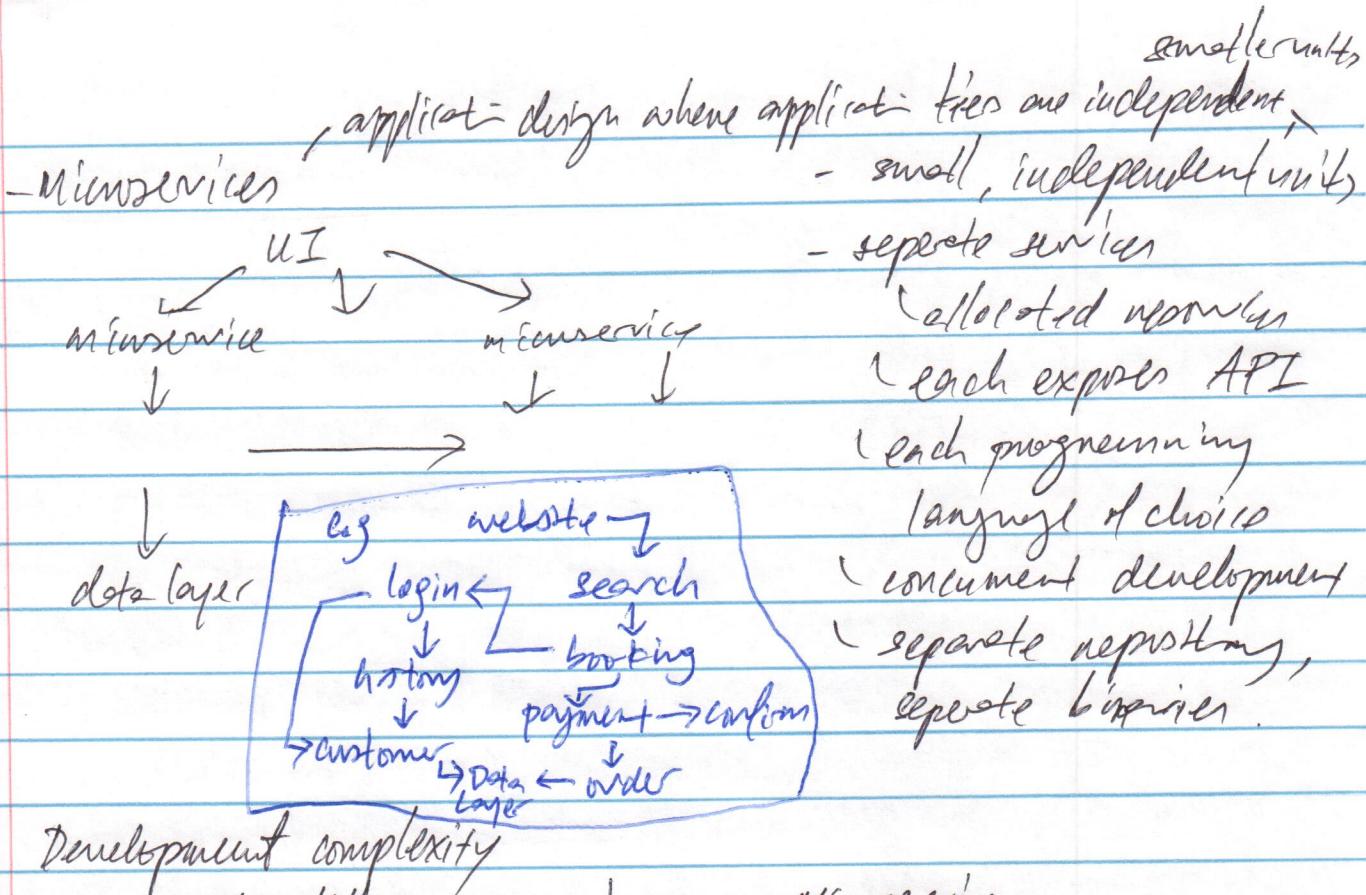
- Requirements
  - stakeholders
  - functionalities
  - end-users
  - input, output processes
  - engineering teams
- Available resources
  - engineering resources
  - budget
  - time frames
  - internal knowledge

### Application architecture

- monolith, microservice
- application tiers
  - UI - HTTP requests + response
  - business logic - services to users
- data layer - access + storage of data layers



- same code
- same repository, share resources, same programming language.  $\rightarrow$  easier



### Development complexity

- Monolith
- 1 language
  - 1 repository
  - sequential

### Microservices

- >1 languages
- >1 repositories
- concurrent

### Scalability

- replicate of entire stack
- overconsumpt. of resources

- replicate of individual functionality
- on-demand resource consumption

### Time to deploy

- 1 delivery pipeline
- entire stack development  
(failure on 1, brings down entire app)

- multiple delivery pipelines
- separate func. deployment
- high velocity at scale

Flexibility

- low

- high

Operational Cost

- low initial cost
- high cost of scale

- high initial cost
- low cost at scale

Reliability

- recovery of entire stack
- low visibility into logs

- recovery of failed component only
- high visibility into logs

Best practices

- health checks
  - status report
    - /healthz or /status
    - return response - 'healthy' or 'unhealthy'
  - metrics
    - statistics for individual services
      - (CPU, memory, network throughput)
    - used + handled numbers
    - /metrics
- logs
  - record operations
  - return response
  - passive logging: \$TOUT, \$TDERR
  - active logging: Direct interface
- tracing
  - measure requested lifecycles
  - application layer implementation

DEBUG,  
INFO  
WARN  
ERROR  
FATAL

CPU I

- CPU + memory - return response - memory usage
- network throughput - " " - 100 MB/s

## Maintenance

- extensibility - microservices
- flexibility - monoliths

## Long-term Applications

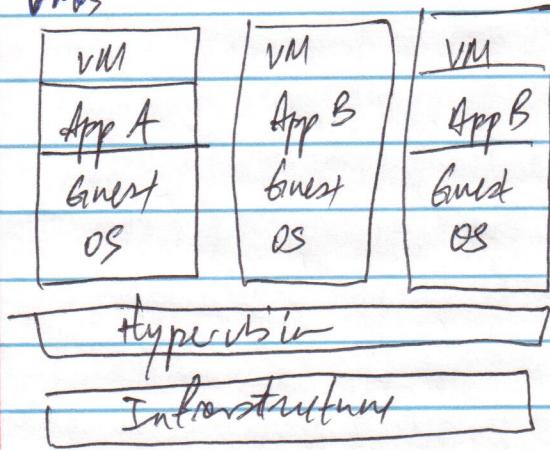
- split functionality
- merge "
- replace "
- stale "

## 2.3 Container Orchestration w/ Kubernetes

### Carries

- declarative application packaging
- container orchestration w/ Kubernetes
- Kubernetes resources
- declarative Kubernetes manifests

## VMs



- standardize .
- multiple VMs on same physical machine .
- replicated OS

- Hypervisor - Manager VM  
- provides virtualization of infrastructure

Source code  
config files  
dependencies

Container

Docker  
Container manager

Docker file  
Docker Image  
Docker Registry

Docker File - set of instructions

- create docker image
- option to package code + dependencies
- layers

- FROM - sets base image

RUN - execute command

COPY & ADD - copy files from host to container

CMD - set default command to execute when container starts.

EXPOSE - expose application port.



## Containers



- Docker - contains management tool
- lightweight
- better usage of resources

Docker image - read-only template

>Create a run-able instance

- docker build [OPTIONS] PATH

-- tag - name of tag

-- file - name dockerfile

- docker run [OPTIONS] IMAGE [COMMAND] [ARG...]

-- detach - run in background

-- publish - expose container port to host

-- it - interactive shell

e.g. docker run -d -p 5000:5000 python-helloworld

- docker logs {container-ID}

Docker registry - store + distribute

- tag - version control

docker tag SOURCE\_IMAGE[:TAG] TARGET\_IMAGE[:TAG]

- e.g. docker tag python-helloworld repo-name/python-helloworld:v1.0.0

repo name  
image name  
version tag

- docker push NAME[:tag]

e.g. docker push repo-name/python-helloworld: v1.0.0

- OCI (Open Container Initiative) compliant images

- standardize images

- Buildpacks - without Dockerfiles, do not expose Docker sockets

- Podman

< Buildah

Dockerfile - FROM python:3.8

LABEL maintainer="Name Here"

COPY . /app

WORKDIR /app

RUN pip install -r requirement.txt

CMD ["python", "app.py"]

go init foldername

docker ~~build~~ build -t python-helloworld .

docker images

docker run -d -p 1000:5000 python-helloworld

docker ps

→ local browser 127.0.0.1:5000

docker tag python-helloworld repo-name/python-helloworld:

? docker push repo-name/python-helloworld: v1.0.0

v1.0.0.

docker login

docker pull repo-name/python-helloworld: v1.0.0

Kubernetes - container orchestrator

- portability - platform agnostic

- scalability - horizontal pod autoscaler (HPA)

- resilience → elasticity

  ↳ automated remediation

  ↳ self-healing

  → Replicated

  → Readiness

  ↳ liveness probes

- service discovery - cluster-level domain name system (DNS)

  → distribute traffic to load balance

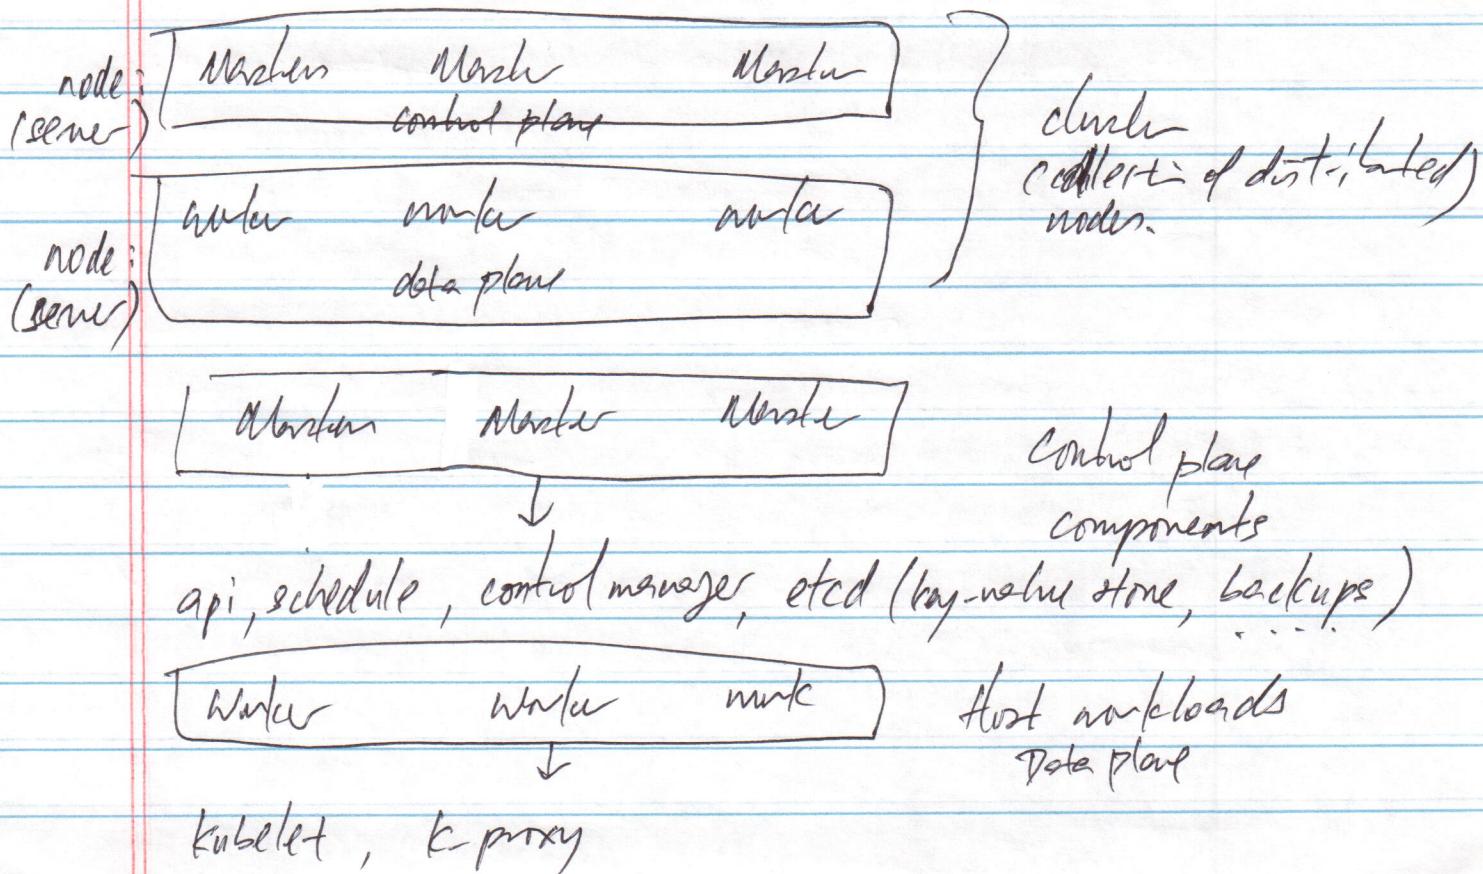
- extensibility - building-blocks principle

  → custom resource definition (CRD)

- operation cost - powerful scheduler

  → only occupies infrastructure it needs

  → cluster autoscaler



kubellet - agent that runs on every node, notifies kube-api-server that this node is part of cluster.

kube-proxy - a network proxy that ensures reachability + accessibility of workloads placed on this node.

kubellet + kube-proxy - installed on all nodes in cluster  
(master and worker nodes)

- keeps kube-api-server up-to-date with list of nodes in cluster + manages connectivity + reachability of workloads.

cluster creation - bootstrap - kubeadm  
- kubespray  
- kops  
- k3s  
- kind  
- minikube } development

k3s - lightweight  
- binary install  
- operational 1-node cluster  
- kubetl - running cluster destroyer

vagrantfile

vagrant up

vagrant status, open virtualbox

vagrant ssh

curl -sPL https://get.k3s.io/k3s.sh  
sudo sh  
kubectl get no

kubeconfig - alien to cluster

- ~/.kube/config

/etc/rancher/k3s/k3s.yaml for k3s

- export KUBECONFIG

- cluster - metadata name, server, certificate-authority

- user - username, authentication - certificate-key

- token

- context - name, cluster, user

- links cluster to user

kubectl cluster-info - get cluster + add-on

" get nodes - get cluster nodes

" get nodes -o wide

" describe node {{ NODE NAME }}  
IP CIDR, Pods, memory

kubernetes - pods - manager applies  
resources

- deployments + replicsets - oversees pods

- services + ingresses - connectivity + reachability to pods

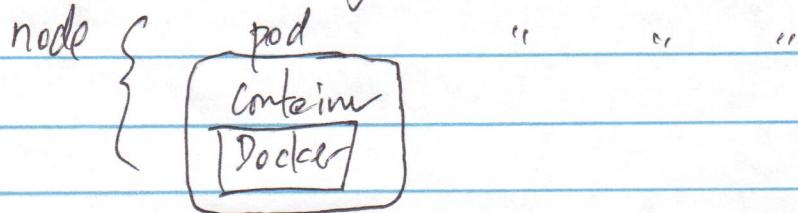
- configmaps + secrets - passes information to pods

- namespaces - separates between multiple applications + their resources

- custom resource definitions (CRD)

extends Kubernetes API to support custom resources.

. Pods - smallest manageable units in Kubernetes cluster.



1 pod : 1 container

1 node ≥ 1 pod

ReplicaSet - ensures desired amount  
of replicas for app is up,  
running.

Deployment → ReplicaSet → Pods

kubectl create deploy NAME --image=image [FLAGS]  
--[COMMAND] [args]

- replicas - set amount of replicas
- namespace - set namespace to run
- port - expose container port

e.g. kubectl create deploy go-helloworld

--image repository/go-helloworld:v1.0.0  
--namespace testspace

Create Pod:

kubectl run NAME --image=name [FLAGS]  
--[COMMAND] [args]

- restart - restart (Always, OnFailure, Never)
- dry-run [None, Client, Server]
- it - open interactive shell to container

e.g. kubectl run -it busybox-test --image=busybox  
--restart=Never

kubectl get deploy  $\Rightarrow$  deployment  
kubectl get rs  $\Rightarrow$  replicates  
" " po  $\Rightarrow$  pods

kubectl edit deploy {{NAME}} -o yaml  
kubectl port-forward {{RESOURCE}} {{PORT}}

Rolling Update - updates pods in a rolling out fashion.  
Recreate - kills all existing pods before new ones are created.

$\hookrightarrow$  minimum down time when upgrading.

Service -

pod 1.2.3.4

pod 1.2.3.5

service;

pod " -6  $\leftarrow$  abstract layer  $\leftarrow$  workload

pod " -7 10.103.44-200

pod " -8

(cluster IP - exposes service to internal cluster IP).

NodePort - exposes service using port exposed on all nodes in cluster.

LoadBalancer - exposes services thru load balancer for public cloud provider.

kubectl expose deploy NAME --port=port [- -target-port=port]  
[FLAGS]

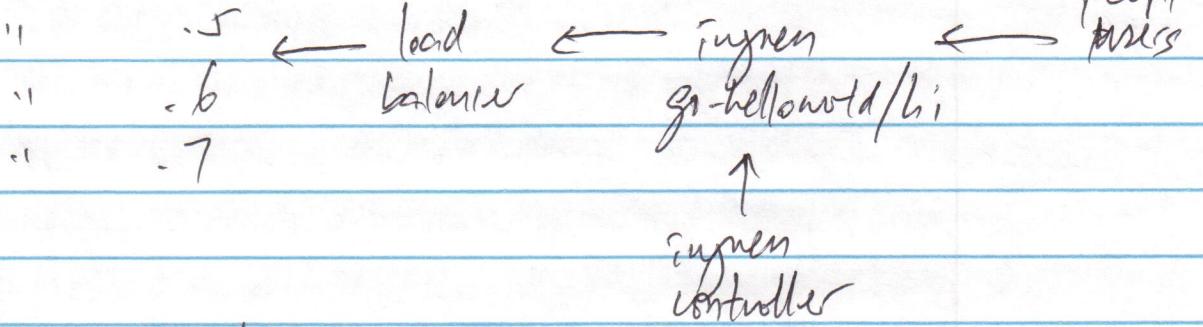
--protocol - network protocol [TCP | UDP | SCTP]

--type - type of service [ClusterIP, NodePort, LoadBalancer]

e.g. kubectl expose deploy go-helloworld --port=8000  
--target-port=6112

Ingress-

pod 1.2.3.4



- express HTTP(S) routes to services within cluster using a load balancer promised by cloud provider.
- set of rules that map HTTP(S) endpoints to services running in cluster.
- Ingress controller keeps rules + load balancer up-to-date

kubectl get svc  $\Rightarrow$  services

kubectl run test-\$RANDOM --namespace=default

--rm -it --image=alpine --sh

wget -qO  $\underbrace{10.105.166.228}_{\text{server IP}}$ :6112

Configmap - non-confidential data

- stored as environmental variables

Config file via a config  
command-line argument

kubectl create configmap NAME [flags]

--from-file - file path to key-value pairs

--from-literal - key-value pair from command-line

e.g. kubectl create configmap test-cm --from-literal =  
key = value

secret - sensitive data

- consumed as environmental variable  
configfile via a volume

• kubectl create secret generic NAME [flags]

--from-file

--from-literal

}

} encrypt by ~~base64~~ base64

Namespace - logical separate between applications

- CPU + memory bound

kubectl create ns NAME

" get po -n NAME

kubectl get cm

" describe cm NAME

" " secrets NAME

" logs RESOURCE/NAME [flags]

" delete "

## Kubernetes manifests

### - improvements

- ↳ line

- ↳ development

- ↳ low learning curve

### - declarations

- ↳ files stored locally

- ↳ predictable

- ↳ high learning curve

## YAML - data serialization standard

- apiVersion {{ API version }}

- kind : {{ Resource }} - labels :

- metadata :

- apps:

- spec: - container: - strategy

- name:

- template:image:

- namespace:

- name:

declarative view - kubectl apply -f deploy.yaml

template:

containers:

name:

ports:

resources:

liveness Probe:

readiness:

readiness Probe:

cpu:

## kubectl create RESOURCE [REQUIRED FLAGS]

- dry-run = client -o yaml

kubectl apply -f manifest.yaml  
" " delete -f "

- deployment (Replicates)

PodSpec - maximum pods - ports

- template probe - pod is running, restart if needed

- service

-

- Control plane
  - if down, applications still work/runs
  - but can't start new workloads, no changes to existing ones
- Data plane
  - still runs if control plane is down

## 2.4 Open Source Platform as a Service (PaaS)

- Covers
- PaaS mechanisms
  - Cloud Foundry
  - Functions as a service

### Cluster management

- sandbox
  - staging
  - production
- } x number of regions
- upgrade, update, manage, configure  
- need a dedicated team in each region.

### PaaS mechanisms

- cloud services
- on-premises
- IaaS (Infrastructure as a service)
- PaaS

### Infrastructure services

- |                                |                  |                         |                |
|--------------------------------|------------------|-------------------------|----------------|
| IaaS                           | - networking     | - OS                    | - applications |
| - AWS, GCP                     | - storage        | PaaS                    |                |
| - on-demand                    | - servers        | - middlewares (API, ..) |                |
| (1b) - self-managed Kubernetes | - virtualization | - runtime (JVM, ..)     |                |
|                                |                  | - data (SQL, MongoDB)   |                |

on-premises - manage all layers on their own  
- large engineering team

PaaS - <sup>PaaS</sup> time efficiency - 3rd party handles management of infrastructure

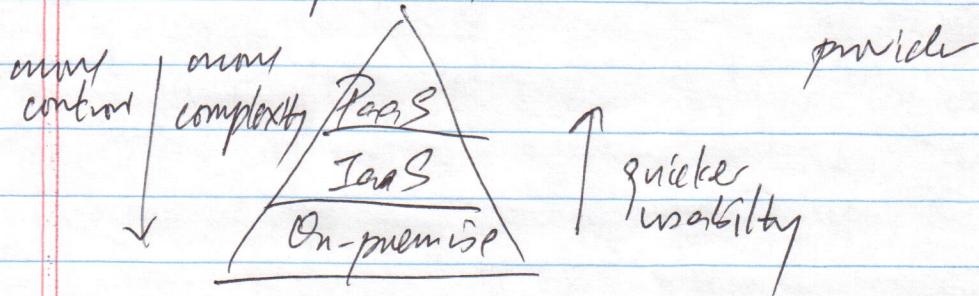
- scalability + high availability
- rich application catalog

- Cons

- vendor lock-in

- data security - data located off-ITP

- operational limitation - limited to skills of provider



Cloud Foundry - open source PaaS

- no vendor lock-in

- handles application lifecycle - build, deployment, execution, scalability

- deployed on top of Kubernetes'

- routing - handles incoming external traffic + route to application

authenticator - manages identities to user accounts

application lifecycle - controls application deployments, monitors status, handles changes

application storage + cache - handles artifacts to application

- DevOps - use service brokers to provision on-demand dependency services for application (e.g. database)
- managing - communicate between Cloud Foundry components.
- metrics + logging -

Kubernetes steps:

1. create OCI compliant image, created with Docker.
2. deploy Kubernetes cluster with a valid ingress controller for routing of requests.
3. deploy an observability stack, includes logs, metrics.
4. create YAML manifests for application deployment.
5. create CI/CD pipeline to push Kubernetes resources to cluster.

Cloud Foundry steps:

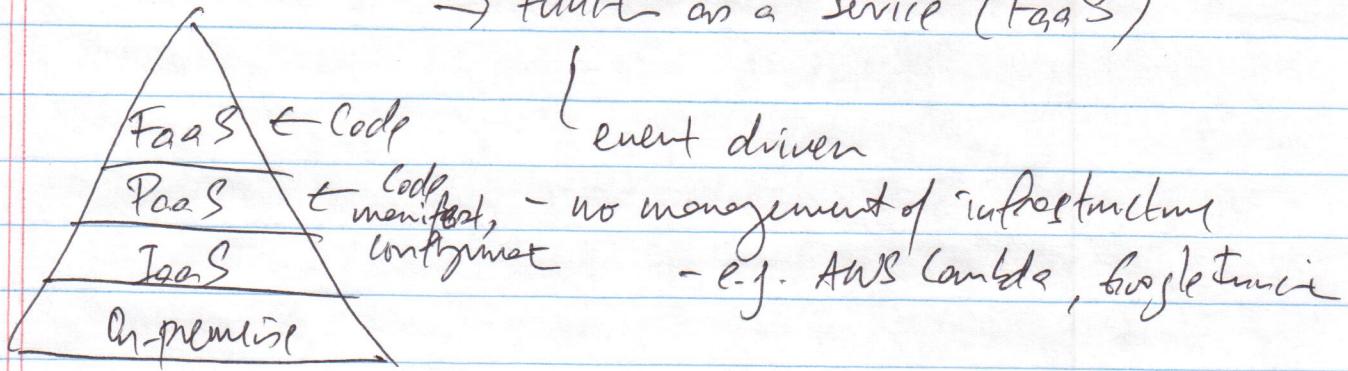
1. write manifest file to provide application deployment parameters.
2. deploy Cloud Foundry or PaaS solution from 3<sup>rd</sup> party.
3. deploy application to Cloud Foundry via CLI or UI
- \* Cloud Foundry will create OCI compliant image by default, provide runtime configuration too.

Cloud Foundry - ~~less~~ greater abstraction (less management of underlying infrastructure.)

- looks like PaaS vendor

Kubernetes - full control over container orchestration, more flexibility in management of applications.

ephemeral services - run 1 to 2 times per day  
 - not cost efficient, billed even not in use  
 → function as a service (FaaS)



## 2.5 CI/CD with Cloud Native Tooling

- Comm - CI/CD
- CI fundamentals
  - CD fundamentals
  - Configuration managers - e.g. Helm

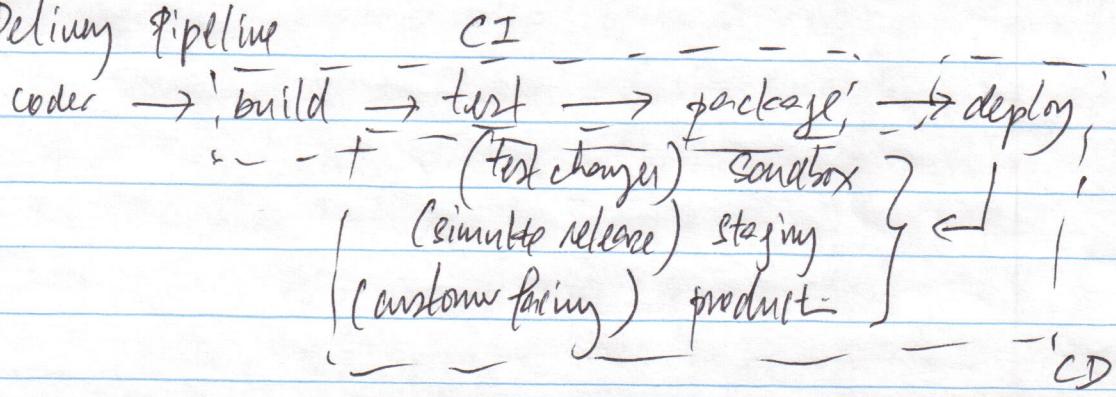
### Cloud Foundry

- simple
- automatically watch new commits
- one click redeploy

### Kubernetes

- manual
- non-scalable
- need to build a delivery pipeline

### Delivery Pipeline



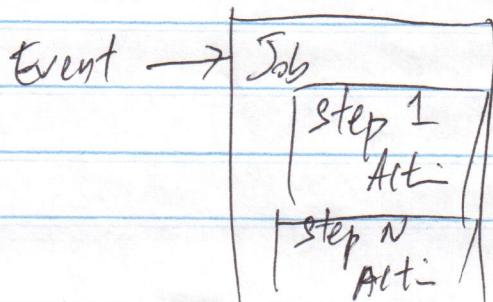
## Advantages of CI/CD

- frequency releases
- less risk - less manual intervention & configurations
- developer productivity - .

## CI

- merge of code
- triggered at each commit
- output: Docker image or binary
- build: 

```
RUN pip install -r requirements.txt  
CMD [ "python", "app.py" ]
```
- test: pytest - functional tests  
pylint - static code analysis
- package: docker build -t python-helloworld  
docker push repo-name/python-helloworld:v1.0.0
- Jenkins, Circle CI, Concourse, Spinnaker, GitHub action
- GitHub action
  - ↳ event-driven - new commits  
- external event  
- scheduled events
  - ↳ multi-language
  - ↳ notification
  - ↳ status badges



• `git hub/workflows/python-version.yaml`

- name : Python version
- on : [push]
- jobs :

check-python-version :

runs-on : ubuntu-latest

steps :

- uses : actions/checkout@v2

- uses : " /setup-python@v2

- run : python-version

e.g. `.github/workflows/docker-build.yaml`

name : Docker build and push

on :

push :

branches : [master]

pull requests

branches : [master]

jobs :

build :

runs-on : ubuntu-latest

steps :

- name : Checkout

- uses : actions/checkout@v2

- name : Set up QEMU

- uses : docker/setup-qemu-action@v1

- name : Set up Docker Buildx

- uses : docker/setup-buildx-action@v1

- name: Login to DockerHub

user: docker/login-action<sup>(2)</sup> v1

with:

username: \${secrets.DOCKERHUB\_USERNAME}

password: \${secrets.DOCKERHUB\_TOKEN}

- name: Build + push

user: docker/build-push-action<sup>(2)</sup> v2

with:

context:

file: ./Dockerfile

platform: linux/amd64

push: true

tags: \${YOUR\_DOCKER\_HUB\_REPOSITORY} //

python-helloworld: latest

## Continuous Delivery (CD)

Push code to

1. sandbox

- new changes tested

2. staging

- release simulated

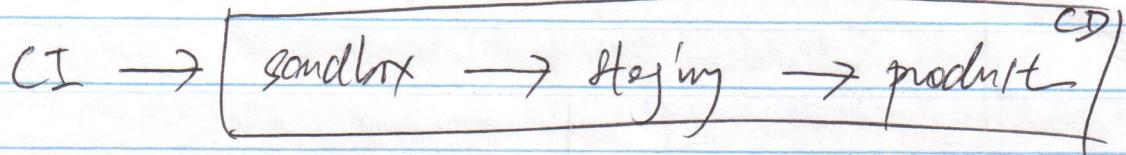
3. product

- customer facing environment.

- regression validation + triggering

}

automated



CI covers deploy stage

- imperative approach

- kubectl create deploy python-helloworld

- --repo-name/python-helloworld:v1.0.0

- declarative approach

- kubectl apply -f deployment.yaml

CI Tools:

- Jenkins

- CircleCI

- Concourse

- Spinnaker

- ArgoCD

ArgoCD

- declarative CD for Kubernetes, follows GitOps.

- configurations stored in manifests.

- monitors new commits in Git repos.

- reconciles changes automatically or on manual trigger.

Advantages

- automatic deployment

- support for multiple config tools

- multi-cluster

- automatic reconciliations

## Angular CD - Resources

### - Projects

- Custom Resource Definition (CRD)
- manage deployment of Kubernetes resources

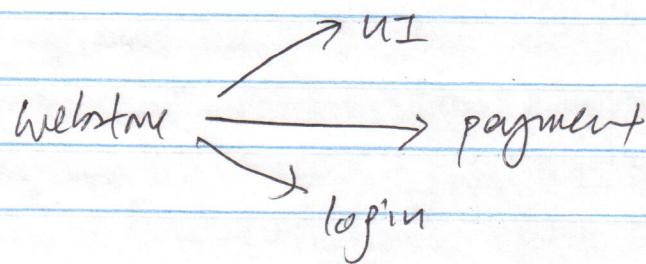
### - Application

apiVersion:

kind:

metadata:

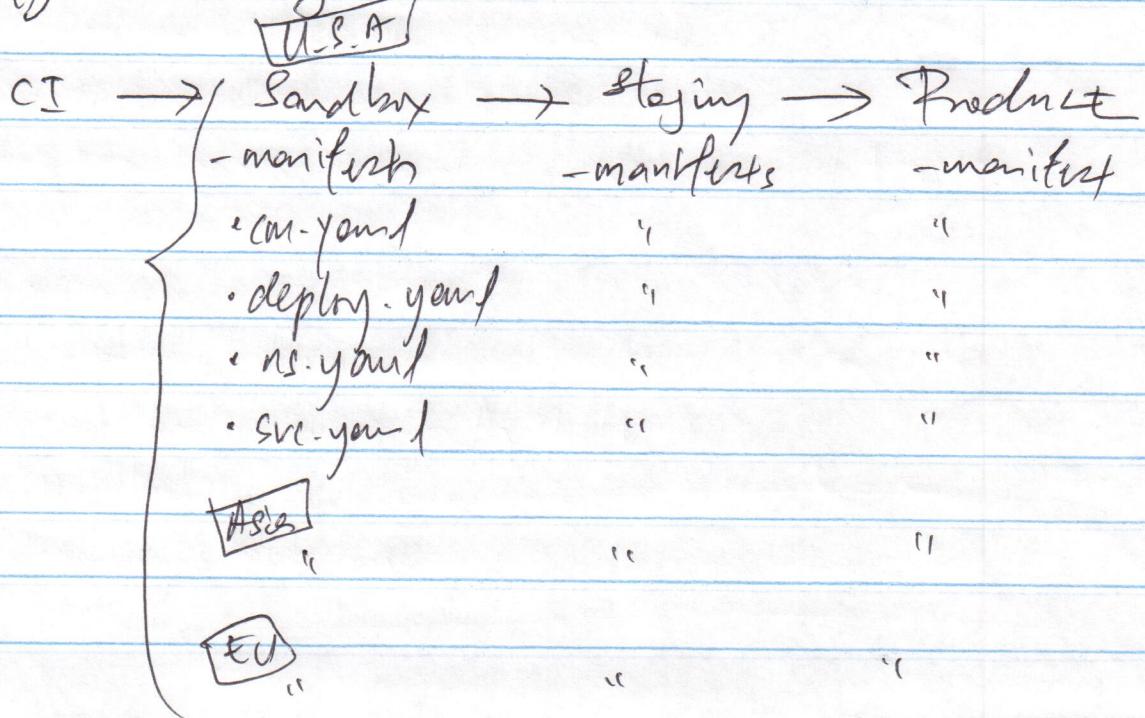
spec:



App of apps:

- group of apps deployed  
+ configured together

## Deployment Manager



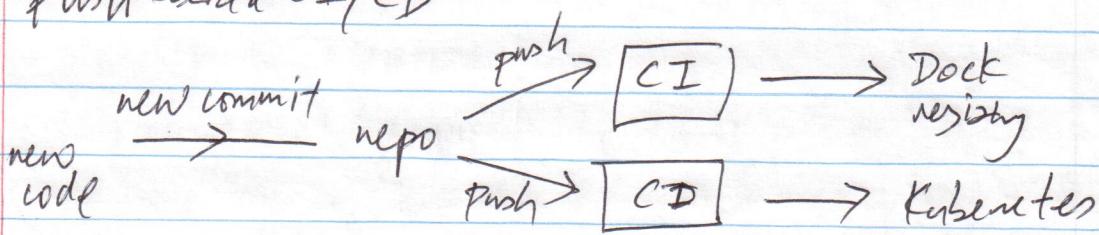
## Tools:

- Helm - templates existing manifests + new files.
- Kustomize - template free, uses a base + multiple overlays.
- Jsonnet - programming language, templating of manifests as JSON files.

## Helm:

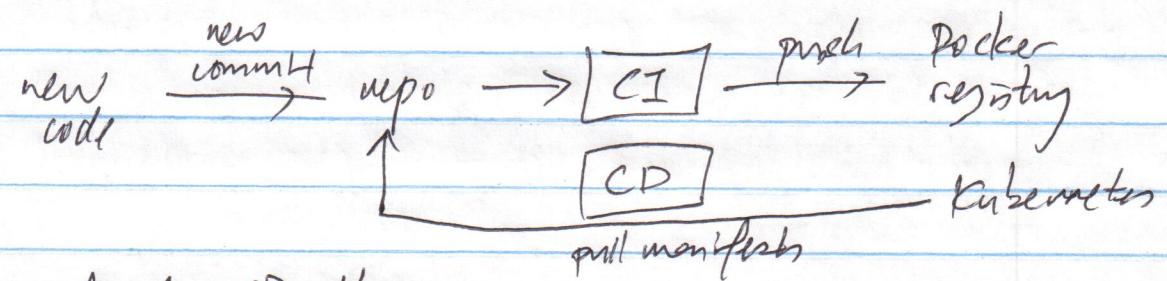
- manages Kubernetes manifest through charts.
- Chart
  - collection of YAML files that describe state of multiple Kubernetes resources.
  - parameterized using GO
  - consists of:
    - chart.yaml
    - templates
      - └ configmap.yaml
      - └ deployment.yaml
      - └ namespaced.yaml
      - └ service.yaml
    - values.yaml

## Push-based CI/CD



e.g. Jenkins, CircleCI

## Pull-based CI/CD



e.g. ~~Argo~~ ArgoCD, Flux.